

Fast Fourier Transforms

Fast Fourier Transforms

- Fourier transforms are useful
- Can we calculate them *efficiently*?

The Fourier transform treats continuous functions defined over infinite domains

$$H(f) = \int_{-\infty}^{\infty} h(t) \exp(-2\pi i f t) dt$$
$$h(t) = \int_{-\infty}^{\infty} H(f) \exp(2\pi i f t) df$$

We can represent a signal over a finite domain with discrete coefficients

If our signal is measured over a period $0 \leq t \leq T$, and if we assume $h(t)$ is periodic, the Fourier transform becomes a Fourier series with components at discrete frequencies $f_n = n/T$.

$$H_n = \int_0^T h(t) \exp(-2\pi i n t / T) dt$$
$$h(t) = \frac{1}{T} \sum_{n=-\infty}^{\infty} H_n \exp(2\pi i n t / T)$$

i.e. an **infinite** number of discrete frequencies can represent a **continuous** periodic signal.

Can sample the signal in the time domain to make the time domain discrete

Consider sampling the signal $h(t)$ at N **uniformly-spaced** points:

$$t_k = k\Delta, k = 0, 1, \dots (N - 1)$$

with $T = N \Delta$

The **sampling rate** is $f_s = 1/\Delta$.


There is a **maximum representable frequency** in such a signal, called the **Nyquist critical frequency** f_c , where we have two samples per cycle:

$$f_c = \frac{1}{2\Delta}$$

Nyquist-Shannon sampling theorem: If $h(t)$ is **band-limited** to frequencies $|f| < f_c$ (i.e. if $H(f) = 0$ for $|f| > f_c$) then $h(t)$ is completely determined by h_n .

Higher frequencies than the Nyquist frequency are aliased because the frequency domain is now periodic

The frequency $f + 2f_c = f + (1/\Delta)$ produces exactly the same samples as f :



Ideally we bandpass filter the signal before sampling to ensure it is bandwidth-limited and then no aliasing can occur.

The Discrete Fourier Transform (DFT) is periodic and discrete in both domains

Given our N samples h_k , we can construct N frequencies which approximate the continuous Fourier transform, with the highest frequency being the critical frequency f_c . It's simplest for now to assume that N is even. We define the **discrete Fourier transform** as

$$H_n = \sum_{k=0}^{N-1} h_k e^{-2\pi i k n / N}$$

which maps N time-domain samples into N frequencies, which are

$$f_n = \frac{n}{N\Delta} = \frac{2n}{N} f_c$$

We now have a discrete signal in the time and frequency domains, with the functions **periodic in both domains**: $h_{k+N} = h_k$ and $H_{n+N} = H_n$.

We can represent negative frequencies using positive indices

The discrete frequencies are $f_n = n/(N\Delta) = 2nf_c/N$, with n running from $n = 0$ to $(N-1)$:

- $n = 0$ is zero frequency (sum of input values)
- $1 \leq n \leq (N/2)$ are positive frequencies, with $(N/2)$ being the highest (Nyquist critical frequency f_c)
- $(N/2) + 1 \leq n \leq (N-1)$ can be thought of as **negative frequencies**: we can subtract $2f_c = 1/\Delta$ from them and they are the same because H_n is periodic

There is an exact inverse

$$h_k = \frac{1}{N} \sum_{n=0}^{N-1} H_n e^{2\pi i k n / N}$$

Note the division by N .

To get the right scaling with respect to the continuous transform we use $H(f_n) \approx \Delta H_n$.

The DFT generalizes to higher dimensions straightforwardly

- If data lives d dimensions with N_i datapoints along dimension $i = 1, \dots, d$

$$F_{\mathbf{n}} = \sum_{\mathbf{j}} f_{\mathbf{j}} e^{-i\eta_{\mathbf{n}} \cdot \mathbf{j}}$$

$\mathbf{j} = (j_1, \dots, j_d)$ with $j_i = 0, \dots, N_i - 1$ and $\eta_{\mathbf{n}} = 2\pi(n_1/N_1, \dots, n_d/N_d)$ $n_i = 0, \dots, N_i - 1$

The Fast Fourier Transform (FFT) is an efficient method for calculating the DFT

- How much computation is involved in a DFT? We can write

$$H_n = \sum_{k=0}^{N-1} W^{nk} h_k$$

where

$$W \equiv \exp(-2\pi i/N)$$

- This looks like a matrix multiplication with a square matrix W whose $N \times N$ elements W_{nk} multiply the vector h_k of length N . This is an $\mathcal{O}(N^2)$ process i.e. its compute time is dominated by a number of complex multiplications proportional to N^2

The FFT is many orders of magnitude faster

- In fact, the FFT algorithm can do the same job in $\mathcal{O}(N \log_2 N)$ operations.
- For $N = 10^6$, $N^2/(N \log_2 N) \approx 50,000$. (50,000s is 14 hours)
- "Discovered" by Danielson & Lanczos (1942), computer discovery by Cooley and Tukey (1965), but the original ideas goes back at least as far as Gauss (1802).

The key to the FFT is writing a DFT of length N as the sum of two DFTs of length $N/2$

Split the DFT into odd and even terms: we can write H_n as

$$\begin{aligned} & \sum_k^{\text{even}} h_k \exp(-2\pi i k n / N) + \sum_k^{\text{odd}} h_k \exp(-2\pi i k n / N) = \\ & \sum_{m=0}^{N/2-1} h_k \exp(-2\pi i (2m) n / N) + \sum_{m=0}^{N/2-1} h_k \exp(-2\pi i (2m+1) n / N) = \\ & \sum_{m=0}^{N/2-1} h_k \exp(-2\pi i m n / (N/2)) + \exp(-2\pi i n / N) \sum_{m=0}^{N/2-1} h_k \exp(-2\pi i m n / (N/2)) \end{aligned}$$

Separate a big FT into 2 smaller FT

If N is a power of 2, we can apply this theorem **recursively**


- Since $2(N/2)^2 = N^2/2$, two half-sized DFTs are $\sim 2\times$ quicker to calculate.
- After $\log_2(N)$ halvings we end up with $N = 1$.
- Highly recommended to use 2^N if efficiency important. Either design your experiment correctly, or **pad with zeros** until your data blocks have length 2^N !
- But fast techniques also exist when N can be factorised.
- Worst case: N is prime, but there exist algorithms to speed up even this case.

FFT Applications

1. Convolution of two signals: FFT each, multiply the FFTs, then inverse FFT back. For example, smooth an image using a Gaussian kernel.
2. **Filtering** a signal – closely related to convolution. We take a signal, FFT it, multiply the FFT by a function, then FFT back, e.g. low- or high-pass filtering.
3. Crystallography
4. Find the power spectrum (PSD) $|H_n|^2$
5. Optics: Fraunhofer (and Fresnel) diffraction, spatial/temporal coherence function
6. Signal processing. e.g. Freeview signals are transmitted using FFT: the data are cut into 8192-piece chunks, ($8192 = 2^{13}$), FFT'ed, transmitted, inverse FFT'ed on reception. In fact, FFTs are central algorithm in Digital Signal Processing (DSP)


To test FFTs in Python, we first generate some fake data

```
import numpy as np
import matplotlib.pyplot as plt
dt=0.01
fftsize=256
t=np.arange(fftsize)*dt
#Generate some fake data at 12 Hz and 34 Hz
y=np.cos(2*np.pi*12*t)+0.5*np.sin(2*np.pi*34*t)
plt.plot(t,y)
```




Negative frequencies come *after* positive frequencies

```
Y=np.fft.fft(y)
# Plot FFT modulus versus array index
plt.subplot(2,1,1); plt.plot(abs(Y))
# Now use the correct frequency coordinates
f=np.fft.fftfreq(fftsize,dt)
plt.subplot(2,1,2); plt.plot(f,abs(Y))
```



Negative frequencies are not always needed


```
plt.subplot(2,1,1); plt.plot(f,Y.real)
plt.subplot(2,1,2); plt.plot(f,Y.imag)
```



Recall that the FT of a real signal has $H(f) = H^*(-f)$

Re-ordering the array makes plots tidier


```
Y2=np.fft.fftshift(Y)
f2=np.fft.fftshift(f)
plt.subplot(2,1,1); plt.plot(f2,Y2.real)
plt.subplot(2,1,2); plt.plot(f2,Y2.imag)
```




In 2-d negative-frequency re-ordering is even more helpful to visualise

Fourier results


```
fig, (ax1, ax2) = plt.subplots(1, 2)
ax1.imshow(smiley, cmap="gray")
ax2.imshow(np.fft.fftshift(smiley), cmap="gray")
```




Signal Processing Example - LIGO



Plotting the data in the Fourier domain shows the noise sources





We can Fourier filter the data and inverse transform to reveal the underlying signal

